

# 操作系统 第三次实验报告

周子锐 2100011032

2023 年 10 月 27 日

完成的 Challenge 为:

Challenge! Modify the JOS kernel monitor so that you can 'continue' execution from the current location (e.g., after the `int3`, if the kernel monitor was invoked via the breakpoint exception), and so that you can single-step one instruction at a time. You will need to understand certain bits of the EFLAGS register in order to implement single-stepping.

## Exercise 1

与 lab2 类似, 根据注释完成代码.

```
////////////////////////////////////
// Make 'envs' point to an array of size 'NENV' of 'struct Env'.
// LAB 3: Your code here.
envs = (struct Env*) boot_alloc(NENV * sizeof(struct Env));
...
////////////////////////////////////
// Map the 'envs' array read-only by the user at linear address UENVS
// (ie. perm = PTE_U | PTE_P).
// Permissions:
// - the new image at UENVS -- kernel R, user R
// - envs itself -- kernel RW, user NONE
// LAB 3: Your code here.
boot_map_region(kern_pgdir, UENVS, PTSIZE, PADDR(envs), PTE_U | PTE_P);
```

## Exercise 2

根据注释完成以下函数.

## void env\_init(void)

初始化 envs 数组中的元素, 注意因为要让链表头为 0 所以应倒序遍历.

```
void
env_init(void)
{
    // Set up envs array
    for (int i = NENV - 1; i >= 0; i--) {
        envs[i].env_id = 0;
        envs[i].env_status = ENV_FREE;
        envs[i].env_link = env_free_list;
        env_free_list = &envs[i];
    }
    // Per-CPU part of the initialization
    env_init_percpu();
}
```

## static int env\_setup\_vm(struct Env \*e)

对每一个 env 新建一个页目录, 可以直接使用 Kernel 的页目录作为模板.

```
static int
env_setup_vm(struct Env *e)
{
    int i;
    struct PageInfo *p = NULL;

    // Allocate a page for the page directory
    if (!(p = page_alloc(ALLOC_ZERO)))
        return -E_NO_MEM;

    p->pp_ref++;
    e->env_pgdir = (pde_t *)page2kva(p);
    memcpy(e->env_pgdir, kern_pgdir, PGSIZE);

    // UVPT maps the env's own page table read-only.
    // Permissions: kernel R, user R
    e->env_pgdir[PDX(UVPT)] = PADDR(e->env_pgdir) | PTE_P | PTE_U;
```

```
    return 0;
}
```

```
static void region_alloc(struct Env *e, void *va, size_t len)
```

为 env 申请一段连续的虚拟地址空间, 并映射到物理地址空间. 逐页申请并映射即可.

```
static void
region_alloc(struct Env *e, void *va, size_t len)
{
    if (!len) {
        return;
    }
    uintptr_t vstart = ROUNDDOWN((uintptr_t)va, PGSIZE);
    uintptr_t vend = ROUNDUP((uintptr_t)va + len, PGSIZE);
    if (vstart > vend) {
        panic("region_alloc: invalid va and len");
    }
    for (; vstart < vend; vstart += PGSIZE) {
        struct PageInfo *pp = page_alloc(ALLOC_ZERO);
        if (!pp) {
            panic("region_alloc: page_alloc failed");
        }
        if (page_insert(e->env_pgdir, pp, (void *)vstart, PTE_U | PTE_W) < 0) {
            panic("region_alloc: page_insert failed");
        }
    }
}
```

```
static void load_icode(struct Env *e, uint8_t *binary)
```

将一个 ELF 格式的二进制文件加载到 env 中, 参考 boot/main.c 实现. 这里可以预先切换成新环境的页目录, 以便于 memcpy 复制连续的 pages.

```
static void
load_icode(struct Env *e, uint8_t *binary)
{
    struct Elf *elfhdr;
    struct Proghdr *ph, *eph;
    elfhdr = (struct Elf *)binary;
```

```
if (elfhdr->e_magic != ELF_MAGIC) {
    panic("load_icode: invalid ELF header");
}
ph = (struct Proghdr *)((uintptr_t)binary + elfhdr->e_phoff);
eph = ph + elfhdr->e_phnum;
lcr3(PADDR(e->env_pgdir));
for (; ph < eph; ph++) {
    if (ph->p_type == ELF_PROG_LOAD) {
        region_alloc(e, (void *)ph->p_va, ph->p_memsz);
        memset((void *)ph->p_va, 0, ph->p_memsz);
        memcpy((void *)ph->p_va, (void *) (binary + ph->p_offset), ph->p_filesz);
    }
}
e->env_tf.tf_eip = elfhdr->e_entry;
lcr3(PADDR(kern_pgdir));

// Now map one page for the program's initial stack
// at virtual address USTACKTOP - PGSIZE.
region_alloc(e, (void *) (USTACKTOP - PGSIZE), PGSIZE);
}
```

**void env\_create(uint8\_t \*binary, enum EnvType type)**

分配一个可用的 Env 来加载一个二进制文件.

```
void
env_create(uint8_t *binary, enum EnvType type)
{
    struct Env *e;
    int ret = env_alloc(&e, 0);
    if (ret < 0) {
        panic("env_create: %e", ret);
    }
    load_icode(e, binary);
    e->env_type = type;
}
```

**void env\_run(struct Env \*e)**

切换到给定环境中运行. 根据注释完成即可.

```

void
env_run(struct Env *e)
{
    if ((curenv) && (curenv->env_status == ENV_RUNNING)) {
        curenv->env_status = ENV_RUNNABLE;
    }
    curenv = e;
    curenv->env_status = ENV_RUNNING;
    curenv->env_runs++;
    lcr3(PADDR(curenv->env_pgdir));

    env_pop_tf(&(curenv->env_tf));
}

```

在实现完上述函数后运行 qemu, 可以发现确实触发了一个 triple fault.

```

npages = 32768
check_page_free_list() succeeded!
check_page_alloc() succeeded!
check_page() succeeded!
check_kern_pgdir() succeeded!
check_page_free_list() succeeded!
check_page_installed_pgdir() succeeded!
[00000000] new env 00001000
EAX=00000000 EBX=00000000 ECX=0000000d EDX=eebfde88
ESI=00000000 EDI=00000000 EBP=eebfde60 ESP=eebfde54
EIP=00800b52 EFL=00000092 [--S-A--] CPL=3 II=0 A20=1 SMM=0 HLT=0
ES =0023 00000000 ffffffff 00cfff300 DPL=3 DS [-WA]
CS =001b 00000000 ffffffff 00cffa00 DPL=3 CS32 [-R-]
SS =0023 00000000 ffffffff 00cfff300 DPL=3 DS [-WA]
DS =0023 00000000 ffffffff 00cfff300 DPL=3 DS [-WA]
FS =0023 00000000 ffffffff 00cfff300 DPL=3 DS [-WA]
GS =0023 00000000 ffffffff 00cfff300 DPL=3 DS [-WA]
LDT=0000 00000000 00000000 00008200 DPL=0 LDT
TR =0028 f0181bc0 00000067 00408900 DPL=0 TSS32-avl
GDT= f011c300 0000002f
IDT= f01813a0 000007ff
CR0=80050033 CR2=00000000 CR3=003bc000 CR4=00000000
DR0=00000000 DR1=00000000 DR2=00000000 DR3=00000000
DR6=ffff0ff0 DR7=00000400
EFER=0000000000000000
Triple fault. Halting for inspection via QEMU monitor.
QEMU: Terminated

```

图 1: Exercise 2 完成后的 qemu 运行信息

同时使用 gdb 调试发现这的确是由 int 指令引起的.

```
(gdb) b *0x800d03
Breakpoint 3 at 0x800d03
(gdb) c
Continuing.

Program received signal SIGTRAP, Trace/breakpoint trap.
=> 0x800b52: int $0x30
0x00800b52 in ?? ()
(gdb) si
=> 0x800b52: int $0x30
0x00800b52 in ?? ()
(gdb) q
A debugging session is active.

Inferior 1 [Remote target] will be killed.
```

图 2: Exercise 2 完成后的 gdb 调试信息

## Exercise 4

在 kern/trap.c 中将生成的中断处理函数用 SETGATE 宏加载到 IDT 中, 需要指出的是, 这里为 Breakpoint 和 Syscall 两个中断设置了特殊的 DPL, 以便于在用户态下使用。

```
void handler_divide();
void handler_debug();
void handler_nmi();
void handler_brkpt();
void handler_oflow();
void handler_bound();
void handler_illop();
void handler_device();
void handler_dblflt();
void handler_tss();
void handler_segnp();
void handler_stack();
void handler_gpflt();
void handler_pgflt();
void handler_fperr();
void handler_align();
void handler_mchk();
void handler_simderr();
void handler_syscall();

void
trap_init(void)
{
    extern struct Segdesc gdt[];
```

```

SETGATE(idt[T_DIVIDE], 0, GD_KT, handler_divide, 0);
SETGATE(idt[T_DEBUG], 0, GD_KT, handler_debug, 0);
SETGATE(idt[T_NMI], 0, GD_KT, handler_nmi, 0);
SETGATE(idt[T_BRKPT], 0, GD_KT, handler_brkpt, 3);
SETGATE(idt[T_OFLOW], 0, GD_KT, handler_oflow, 0);
SETGATE(idt[T_BOUND], 0, GD_KT, handler_bound, 0);
SETGATE(idt[T_ILLOP], 0, GD_KT, handler_illop, 0);
SETGATE(idt[T_DEVICE], 0, GD_KT, handler_device, 0);
SETGATE(idt[T_DBLFLT], 0, GD_KT, handler_dblflt, 0);
SETGATE(idt[T_TSS], 0, GD_KT, handler_tss, 0);
SETGATE(idt[T_SEGNP], 0, GD_KT, handler_segnp, 0);
SETGATE(idt[T_STACK], 0, GD_KT, handler_stack, 0);
SETGATE(idt[T_GPFLT], 0, GD_KT, handler_gpflt, 0);
SETGATE(idt[T_PGFLT], 0, GD_KT, handler_pgflt, 0);
SETGATE(idt[T_FPERR], 0, GD_KT, handler_fperr, 0);
SETGATE(idt[T_ALIGN], 0, GD_KT, handler_align, 0);
SETGATE(idt[T_MCHK], 0, GD_KT, handler_mchk, 0);
SETGATE(idt[T_SIMDERR], 0, GD_KT, handler_simderr, 0);
SETGATE(idt[T_SYSCALL], 0, GD_KT, handler_syscall, 3);

// Per-CPU setup
trap_init_percpu();
}

```

接下来是 kern/trapentry.S, 首先使用给出的宏为各类中断生成入口:

```

TRAPHANDLER_NOEC(handler_divide, T_DIVIDE)
TRAPHANDLER_NOEC(handler_debug, T_DEBUG)
TRAPHANDLER_NOEC(handler_nmi, T_NMI)
TRAPHANDLER_NOEC(handler_brkpt, T_BRKPT)
TRAPHANDLER_NOEC(handler_oflow, T_OFLOW)
TRAPHANDLER_NOEC(handler_bound, T_BOUND)
TRAPHANDLER_NOEC(handler_illop, T_ILLOP)
TRAPHANDLER_NOEC(handler_device, T_DEVICE)
TRAPHANDLER(handler_dblflt, T_DBLFLT)
TRAPHANDLER(handler_tss, T_TSS)
TRAPHANDLER(handler_segnp, T_SEGNP)
TRAPHANDLER(handler_stack, T_STACK)
TRAPHANDLER(handler_gpflt, T_GPFLT)
TRAPHANDLER(handler_pgflt, T_PGFLT)

```

```

TRAPHANDLER_NOEC(handler_fperr, T_FPERR)
TRAPHANDLER(handler_align, T_ALIGN)
TRAPHANDLER_NOEC(handler_mchk, T_MCHK)
TRAPHANDLER_NOEC(handler_simderr, T_SIMDERR)
TRAPHANDLER_NOEC(handler_syscall, T_SYSCALL)

```

接着是 `_alltraps` 函数, 这是为 `TRAPHANDLER` 提供的跳转函数. 对于第一条要求, 我们注意到 `kern/trap.h` 中的定义

```

struct Trapframe {
    struct PushRegs tf_regs;
    uint16_t tf_es;
    uint16_t tf_padding1;
    uint16_t tf_ds;
    uint16_t tf_padding2;
    uint32_t tf_trapno;
    /* below here defined by x86 hardware */
    uint32_t tf_err;
    uintptr_t tf_eip;
    uint16_t tf_cs;
    uint16_t tf_padding3;
    uint32_t tf_eflags;
    /* below here only when crossing rings, such as from user to kernel */
    uintptr_t tf_esp;
    uint16_t tf_ss;
    uint16_t tf_padding4;
} __attribute__((packed));

```

同时注意到在执行 `_alltraps` 之前已经将 `tf_trapno` 及其之后的信息倒序压入到栈中了, 所以只需要补充其上方的信息即可.

```

.global _alltraps
.type _alltraps, @function
.align 2
_alltraps:
    pushw $0
    pushw %ds
    pushw $0
    pushw %es
    pushal

    movl $GD_KD, %eax

```



```
movw %ax, %ds
movw %ax, %es
pushl %esp
call trap
```

## Question 1

注意到在运行这些处理程序之前, 部分中断还会额外将错误码压入栈中, 这是使用统一的中断处理程序所不能实现的.

## Question 2

根据上面对 IDT 的设置, `int 14` 这一异常在用户态无法申请, 因此会触发第 13 号异常即 `general protection exception`. 如果操作系统允许用户处罚这一缺页异常, 会使得用户能更加方便地操作虚拟内存, 对程序的安全性造成影响.

## Exercise 5&6

直接加上对应条件判断即可.

```
static void
trap_dispatch(struct Trapframe *tf)
{
    // Handle processor exceptions.
    if (tf->tf_trapno == T_PGFLT) {
        page_fault_handler(tf);
        return;
    } else if (tf->tf_trapno == T_BRKPT) {
        monitor(tf);
        return;
    }

    // Unexpected trap: The user process or the kernel has a bug.
    ...
}
```

## Challenge

这个 Challenge 需要我们在接受到 breakpoint 异常的时候能够继续执行, 并且能够单步执行.

通过查阅[这里](#)可以知道 EFLGAS 的 TF 位可以控制程序的单步执行. 因此可以根据如下的代码在 kern/monitor.c 中添加对 continue 和 stepi 命令的支持.

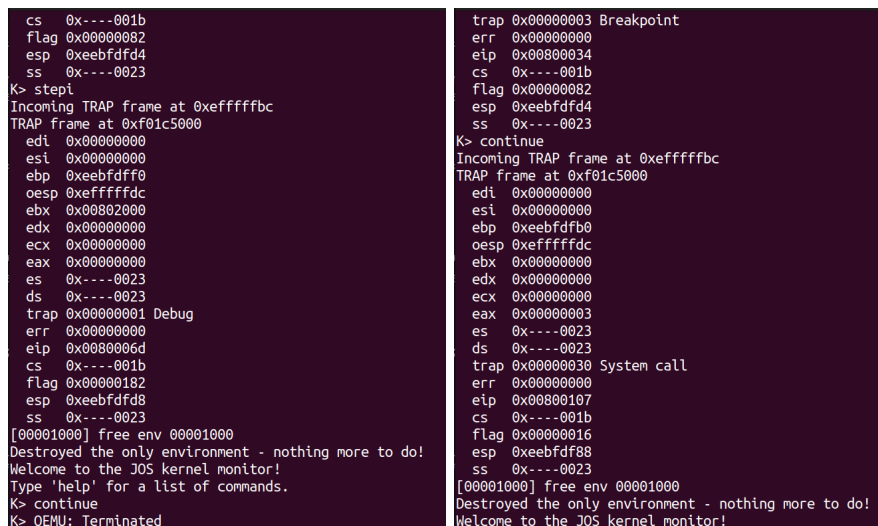
```
int
mon_continue(int argc, char **argv, struct Trapframe *tf) {
    if (!(tf && (tf->tf_trapno == T_DEBUG || tf->tf_trapno == T_BRKPT) &&
        ((tf->tf_cs & 3) == 3)))
        return 0;

    tf->tf_eflags &= ~FL_TF;
    return -1;
}

int
mon_stepi(int argc, char **argv, struct Trapframe *tf) {
    if (!(tf && (tf->tf_trapno == T_DEBUG || tf->tf_trapno == T_BRKPT) &&
        ((tf->tf_cs & 3) == 3)))
        return 0;

    tf->tf_eflags |= FL_TF;
    return -1;
}
```

在 kern/init.c 中修改执行程序为 breakpoint 进行测试得到如下结果.



```
cs 0x----001b
flag 0x00000082
esp 0xeebdfd4
ss 0x----0023
K> stepi
Incoming TRAP frame at 0xeffffbc
TRAP frame at 0xf01c5000
edi 0x00000000
esi 0x00000000
ebp 0xeebdfdf0
oesp 0xefffffdc
ebx 0x00802000
edx 0x00000000
ecx 0x00000000
eax 0x00000000
es 0x----0023
ds 0x----0023
trap 0x00000001 Debug
err 0x00000000
eip 0x0080006d
cs 0x----001b
flag 0x00000182
esp 0xeebdfd8
ss 0x----0023
[00001000] free env 00001000
Destroyed the only environment - nothing more to do!
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> continue
K> QEMU: Terminated

trap 0x00000003 Breakpoint
err 0x00000000
eip 0x00800034
cs 0x----001b
flag 0x00000082
esp 0xeebdfd4
ss 0x----0023
K> continue
Incoming TRAP frame at 0xeffffbc
TRAP frame at 0xf01c5000
edi 0x00000000
esi 0x00000000
ebp 0xeebdfdb0
oesp 0xefffffdc
ebx 0x00000000
edx 0x00000000
ecx 0x00000000
eax 0x00000003
es 0x----0023
ds 0x----0023
trap 0x00000030 System call
err 0x00000000
eip 0x00800107
cs 0x----001b
flag 0x00000016
esp 0xeebdf88
ss 0x----0023
[00001000] free env 00001000
Destroyed the only environment - nothing more to do!
Welcome to the JOS kernel monitor!
```

图 3: 运行 stepi 命令和 continue 命令的结果

### Question 3

因为与其他常规的中断或异常不同, breakpoint 异常是由用户发起的, 当 IDT 中设置其的 DPL 为 0 时, 用户态下若尝试触发此异常则会触发 general protection fault. 因此为了成功触发 breakpoint 需要在加载 IDT 设置其对应项的 DPL 为 3.

### Question 4

这样做是为了防止用户能够随意操作只有操作系统才能操作的资源, 从而实现对 kernel 的攻击.

### Exercise 7

在 kern/trap.c/trap\_dispatch() 中添加对 syscall 的处理.

```
...
else if (tf->tf_trapno == T_SYSCALL) {
    tf->tf_regs.reg_eax = syscall(
        tf->tf_regs.reg_eax,
        tf->tf_regs.reg_edx,
        tf->tf_regs.reg_ecx,
        tf->tf_regs.reg_ebx,
        tf->tf_regs.reg_edi,
        tf->tf_regs.reg_esi
    );
    return;
}
...
```

在 kern/syscall.c 中实现对不同类别的 syscall 的 dispatch, 各个具体 syscall 的 API 参见 lib/syscall.c.

```
int32_t
syscall(uint32_t syscallno, uint32_t a1, uint32_t a2, uint32_t a3, uint32_t a4, uint32_t a5)
{
    // Call the function corresponding to the 'syscallno' parameter.
    // Return any appropriate return value.
```

```
//panic("syscall not implemented");

switch (syscallno) {
case SYS_cputs:
    sys_cputs((const char *)a1, (size_t)a2);
    return 0;
case SYS_cgetc:
    return sys_cgetc();
case SYS_getenvid:
    return sys_getenvid();
case SYS_env_destroy:
    return sys_env_destroy((envid_t)a1);
default:
    return -E_INVALID;
}
}
```

## Exercise 8

在 lib/libmain.c 中加上这一句话即可

```
thisenv = envs + ENVX(sys_getenvid());
```

## Exercise 9

对 kern/pmap.c 中的 user\_mem\_check() 函数, 可以仿照 Exercise 2 中的 region\_alloc() 函数, 逐页判断即可.

```
int
user_mem_check(struct Env *env, const void *va, size_t len, int perm)
{
    if (len == 0) {
        return 0;
    }
    uintptr_t vstart = ROUNDDOWN((uintptr_t)va, PGSIZE);
    uintptr_t vend = ROUNDUP((uintptr_t)va + len, PGSIZE);
    if (vend > ULIM) {
        user_mem_check_addr = MAX(vstart, ULIM);
        return -E_FAULT;
    }
}
```

```

    }
    perm |= PTE_P;
    for (; vstart < vend; vstart += PGSIZE) {
        pte_t *pte = pgdir_walk(env->env_pgdir, (void *)vstart, 0);
        if ((!pte) || ((*pte) & perm) != perm) {
            user_mem_check_addr = MAX(vstart, (uintptr_t)va);
            return -E_FAULT;
        }
    }
    return 0;
}

```

在 kern/syscall.c 中, 目前只有 SYSY\_puts 涉及到访存相关的内容, 为其补充检查即可.

```

static void
sys_cputs(const char *s, size_t len)
{
    // Check that the user has permission to read memory [s, s+len).
    // Destroy the environment if not.
    user_mem_assert(curenv, s, len, PTE_U);

    // Print the string supplied by the user.
    cprintf("%.*s", len, s);
}

```

在 kern/kdebug.c 中, 补充对 usd, stabs, stabstr 的检查即可.

```

int
debuginfo_eip(uintptr_t addr, struct Eipdebuginfo *info)
{
    ...
    // Make sure this memory is valid.
    // Return -1 if it is not. Hint: Call user_mem_check.
    if (user_mem_check(curenv, usd, sizeof(struct UserStabData), PTE_U) < 0)
        return -1;
    ...
    // Make sure the STABS and string table memory is valid.
    if (user_mem_check(curenv, stabs, (uintptr_t)stab_end - (uintptr_t)stabs, PTE_U) < 0)
        return -1;
    if (user_mem_check(curenv, stabstr, stabstr_end - stabstr, PTE_U) < 0)
        return -1;
    ...
}

```

```
}
```

## 测试结果

```
+ ld boot/boot
boot block is 396 bytes (max 510)
+ mk obj/kern/kernel.img
make[1]: Leaving directory '/home/encodetalker/oslabs/lab'
divzero: OK (1.2s)
softint: OK (1.1s)
badsegment: OK (1.0s)
Part A score: 30/30

faultread: OK (1.1s)
faultreadkernel: OK (1.0s)
faultwrite: OK (1.0s)
faultwritekernel: OK (1.0s)
breakpoint: OK (1.0s)
testbss: OK (1.1s)
hello: OK (1.0s)
buggyhello: OK (1.0s)
buggyhello2: OK (1.1s)
evilhello: OK (1.0s)
Part B score: 50/50

Score: 80/80
```

图 4: 评测结果

This completes the lab.